

SpectralStream

A Unified Framework for Neural Network Compression
via Spectral Methods, Hyperdimensional Computing,
and Physics-Inspired Attention

Technical Whitepaper — Version 1.0

SpectralStream Research Group

<https://github.com/anomalyco/SpectralStream>

July 1, 2026

Abstract

We present **SpectralStream**, a pure-Python framework for neural network compression and efficient inference that achieves extreme compression ratios (up to 5000:1) through the synergistic integration of six mathematical traditions: spectral analysis, hyperdimensional computing (HDC), Vlasov mean-field theory from plasma physics, tensor networks from quantum many-body physics, compressed sensing, and information theory. SpectralStream replaces traditional quantization-first approaches with a *compression-first* philosophy that decomposes weight matrices and attention patterns into structured low-entropy representations before optional quantization. The framework encompasses an 80+ method compression intelligence engine with adaptive per-tensor method selection, a production-grade CPU inference pipeline with speculative decoding, a unified binary format (SSF v2/v3) with cryptographic integrity verification, a hierarchical KV cache with 13 eviction policies and 30+ compression methods, and a professional certificate system for quality assurance. All operations are vectorized via NumPy with zero C++ extensions, demonstrating that pure-Python can achieve production-grade neural network compression at scale.

1 Introduction

Large language models (LLMs) have grown to hundreds of billions of parameters, with inference requiring terabytes of memory and significant computational resources. Existing compression approaches—quantization

(GPTQ [9], AWQ [10], GGML [14]), pruning (SparseGPT [11], Wanda [12]), and distillation [13]—operate primarily in the spatial (weight-value) domain and typically achieve 2–8× compression at the cost of measurable quality degradation.

SpectralStream takes a fundamentally different approach: instead of quantizing weights directly, it first *transforms* them into spectral, structural, or tensor-network representations where information is naturally concentrated in far fewer coefficients. The central insight is that neural network weight matrices and attention patterns are not unstructured noise—they exhibit highly structured frequency-domain, low-rank, and low-entanglement properties that can be exploited for compression ratios orders of magnitude beyond conventional quantization.

The core design principles of SpectralStream are:

- Compression-first, quantize-last:** Spectral methods (DCT, FWHT, wavelets), tensor decompositions (SVD, TT, MERA), and structural methods are tried *before* quantization. Quantization is a last resort fallback (Tier 5).
- Adaptive per-tensor intelligence:** The `CompressionIntelligenceEngine` profiles each tensor across 40+ metrics (spectral entropy, effective rank, structural scores, etc.) and selects from 80+ methods via tiered scoring, historical learning, and performance prediction.
- Physics-inspired computation:** $O(n)$ mean-field attention via the Vlasov–Poisson system replaces quadratic $O(n^2)$ self-

attention; hyperdimensional computing enables $O(1)$ speculative decoding.

4. **Pure-Python with NumPy vectorization:** All operations use vectorized NumPy—no C++ extensions, no CUDA. This dramatically simplifies deployment, auditing, and modification.

This whitepaper provides a comprehensive technical description of every component of SpectralStream, organized as follows. Section 2 details the mathematical foundations across six domains. Section 3 describes the compression intelligence engine architecture. Section 4 covers the SSF binary format. Section 5 presents the hierarchical KV cache system. Section 6 details the inference pipeline. Section 7 describes the certificate and validation system. Section 8 presents performance characteristics, and Section 9 concludes.

2 Mathematical Foundations

SpectralStream integrates six mathematical traditions into a unified compression and inference framework.

2.1 Spectral & Harmonic Analysis

The *spectral* core of SpectralStream exploits the observation that neural network weight matrices exhibit strong energy compaction in frequency-domain representations.

Discrete Cosine Transform (DCT). The Type-II DCT is the primary spectral transform, chosen for its near-optimal energy compaction for natural signals and its real-valued (non-complex) output. The 1D DCT-II matrix is defined element-wise as:

$$\text{DCT}(n)_{k,j} = \begin{cases} \frac{1}{\sqrt{n}}, & k = 0, \\ \sqrt{\frac{2}{n}} \cos\left(\frac{\pi k(j+0.5)}{n}\right), & k > 0. \end{cases} \quad (1)$$

The 2D DCT is separable: $\mathbf{Y} = \mathbf{C}_n \mathbf{X} \mathbf{C}_m^\top$, enabling efficient block-wise application. This is the direct mathematical heritage of JPEG image compression, adapted to neural network weight compression [1].

Fast Walsh–Hadamard Transform (FWHT). For Hadamard-based quantization and rotation, the in-place butterfly FWHT is used:

$$\text{FWHT}(\mathbf{x})_i = \sum_{j=0}^{N-1} (-1)^{\langle i,j \rangle} x_j, \quad (2)$$

where $\langle i, j \rangle$ is the bitwise dot product. FWHT is self-inverse (up to scaling) and $O(N \log N)$.

Wavelet Transforms. Multi-resolution analysis via Haar (lifting scheme) and Daubechies-4 wavelets enables scale-dependent compression, where low-frequency approximation coefficients are preserved and high-frequency detail coefficients are thresholded.

Number Theoretic Transform (NTT). For lossless integer-domain convolution, NTT uses primes of the form $p = c \cdot 2^k + 1$ with automatic primitive root detection, enabling exact convolution without floating-point error.

Spectral Analysis Tools. The spectral analyzer computes: *spectral entropy* $H_s = -\sum p_i \log_2 p_i / \log_2 N$ (normalized Shannon entropy of the power spectrum), *energy concentration* (fraction of total energy in top- k frequency bins), *effective rank* $\rho = \exp(-\sum \tilde{\sigma}_i \log \tilde{\sigma}_i)$, and *band-limiting* via low-pass Fourier filtering.

2.2 Hyperdimensional Computing

Hyperdimensional computing (HDC) [2] operates in a high-dimensional vector space (typically 10,000 dimensions) where random vectors are nearly orthogonal with high probability.

Holographic Reduced Representations (HRR). SpectralStream implements the HRR framework of Plate [3]:

$$\mathbf{a} \circledast \mathbf{b} = \mathcal{F}^{-1}\{\mathcal{F}(\mathbf{a}) \odot \mathcal{F}(\mathbf{b})\}, \quad (3)$$

$$\mathbf{a} \# \mathbf{b} = \mathcal{F}^{-1}\{\mathcal{F}(\mathbf{a}) \odot \overline{\mathcal{F}(\mathbf{b})}\}, \quad (4)$$

$$\mathbf{a} \oplus \mathbf{b} = \mathbf{a} + \mathbf{b}, \quad (5)$$

where $\circledast$ denotes circular convolution (binding), $\#$ denotes circular correlation (unbinding), and $\oplus$ denotes element-wise addition (bundling).

Binding creates a vector dissimilar to both operands (for role–filler associations), while bundling maintains approximate similarity to all components (for set formation). The operations satisfy approximate inversion: $(\mathbf{a} \circledast \mathbf{b}) \# \mathbf{a} \approx \mathbf{b}$.

HDC for Speculative Decoding. The `HDCDraftEngine` uses 10,000-dimensional hyper-vectors for ultra-fast draft generation:

- Each vocabulary token maps to a sparse hypervector (5% density, random ± 1 entries).
- N-gram sequences (order 2–6) are bound via circular convolution, creating context-specific vectors.
- Multi-prototype storage maintains up to 8 prototype HD vectors per context length, enabling diverse candidate generation.
- Locality-Sensitive Hashing (LSH) tables enable $O(1)$ nearest-prototype lookup, independent of vocabulary size.

This achieves $O(1)$ next-token prediction—constant time regardless of vocabulary size—fundamentally different from the $O(|V|)$ softmax bottleneck.

2.3 Vlasov Mean-Field Attention

The most novel mathematical contribution is the replacement of $O(n^2)$ quadratic attention with $O(n)$ mean-field attention modeled on the Vlasov–Poisson system from plasma physics [4].

Physical Analogy. Tokens are treated as charged particles in a plasma, evolving according to:

$$\frac{\partial f}{\partial t} + \mathbf{v} \cdot \nabla_{\mathbf{x}} f - \nabla_{\mathbf{x}} \Phi \cdot \nabla_{\mathbf{v}} f = 0, \quad \nabla^2 \Phi = -\rho, \quad (6)$$

where $f(\mathbf{x}, \mathbf{v}, t)$ is the phase-space distribution of tokens, $\rho(\mathbf{x}) = \int f d\mathbf{v}$ is the charge density (token concentration), and $\Phi(\mathbf{x})$ is the mean-field potential (attention field). Interactions are mediated by the Yukawa-screened potential:

$$V(r) = \frac{4\pi}{k^2 + \mu^2} \iff V(r) = \frac{Ae^{-r/\mu}}{r + \varepsilon}, \quad (7)$$

where μ is the Debye screening length controlling interaction range.

Algorithm. The mean-field attention proceeds in $O(n)$ steps:

1. **Charge Deposition** ($O(n)$): Token keys project onto a spectral grid of size $n_{\text{grid}} \ll n$ via cloud-in-cell interpolation: $\rho[j] += w_j \|\mathbf{k}_i\|^2$.
2. **Poisson Solve** ($O(n_{\text{grid}} \log n_{\text{grid}})$): $\mathcal{F}(\rho) \rightarrow \tilde{V}(\mathbf{k}) \hat{\rho}(\mathbf{k}) \rightarrow \mathcal{F}^{-1}(\Phi)$.
3. **Potential Interpolation** ($O(n)$): Grid potential interpolated back to token positions.
4. **Output Weighting** ($O(n)$): $\mathbf{o}_i = \mathbf{v}_i \cdot \text{softmax}(-\Phi_i/T)$.

Variants. Several extensions are implemented: `VlasovFlashAttention` (tiled, $O(b \cdot d_h + n_{\text{grid}})$ per tile), `GyrokineticAttention` (DCT-based splitting into slow/semantic and fast/detail components with quadratic full attention for slow tokens only), `SymplecticAttentionIntegrator` (Hamiltonian leapfrog integration preserving phase-space volume), `HelmholtzDecomposition` (irrotational/solenoidal splitting), `TurbulentCascadeAttention` (Kolmogorov energy cascade across scales), and `AdaptiveDebyeAttention` (learnable screening length per head).

2.4 Tensor Networks

Tensor networks from quantum many-body physics provide exponentially efficient representations of high-dimensional tensors by exploiting low entanglement structure [5].

Matrix Product States / Tensor Trains.

An N -dimensional tensor $\mathcal{T} \in \mathbb{R}^{d_1 \times \dots \times d_N}$ is decomposed into a chain of 3-core tensors via sequential SVD:

$$\mathcal{T}_{i_1 \dots i_N} \approx \sum_{\alpha_1 \dots \alpha_{N-1}} A_{i_1, \alpha_1}^{(1)} A_{\alpha_1, i_2, \alpha_2}^{(2)} \dots A_{\alpha_{N-1}, i_N}^{(N)}, \quad (8)$$

with bond dimensions χ_k controlling compression fidelity. For a 2D weight matrix, this reduces to the standard truncated SVD: $\mathbf{W} \approx \mathbf{U}_k \Sigma_k \mathbf{V}_k^T$.

MERA. The Multi-scale Entanglement Renormalization Ansatz [6] provides a hierarchical binary tree structure with isometries ($\chi \times 2 \times \chi$) and disentanglers ($\chi \times \chi \times \chi \times \chi$) achieving $O(\chi^3)$ parameters independent of N —matching ADNTN results at $2000\text{--}77,000\times$ compression.

Additional Formats. Tensor Ring (periodic boundary conditions), iPEPS (2D grid), Butterfly (recursive Kronecker products), and Einsort (adaptive index reordering) provide specialized representations for different tensor structures.

2.5 Compressed Sensing & Information Theory

FISTA-based L1-minimization [7] enables recovery of weight matrices from random projections:

$$\min_{\mathbf{x}} \frac{1}{2} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2 + \lambda \|\mathbf{x}\|_1, \quad (9)$$

exploiting that weights are approximately sparse in DCT/wavelet bases. The information-theoretic pipeline terminates with entropy coding (ANS [8], arithmetic, Huffman, range coding) achieving near-optimal compression given the quantized coefficient distribution.

3 Compression Intelligence Engine

The `CompressionIntelligenceEngine` is the central orchestrator, implementing a five-stage pipeline: Profile $\rightarrow$ Allocate $\rightarrow$ Select $\rightarrow$ Compress $\rightarrow$ Validate.

3.1 Profiling

Each tensor is profiled across five analysis dimensions (40+ metrics total):

- **Statistical:** mean, std, min/max, kurtosis, skewness, dynamic range, outlier ratio.
- **Spectral:** effective rank, spectral decay rate, energy concentration, spectral entropy.
- **Structural (2D):** Toeplitz score (diagonal constancy), circulant score (column-shift correlation), block-diagonal score (on-diagonal energy ratio), hierarchical score (quadrant correlation).

- **Information:** entropy rate (16-state Markov), mutual information (cross-block correlation), Kolmogorov complexity (LZ dictionary ratio).
- **Sparsity:** N:M structured (2:4, 4:8), block sparsity, unstructured sparsity.

Tensor type is classified by name (embedding, attention Q/K/V/O, FFN gate/up/down, etc.) and loaded with sensitivity weights ranging from 1.0 (embeddings, attention Q/O) to 0.50 (layer norms). Gradient-based sensitivity from calibration Fisher information can supplement name-based sensitivity.

3.2 Error Budget Allocation

The `ErrorBudgetAllocator` distributes error across tensors by sensitivity:

$$\varepsilon_i = \min \left(\varepsilon_{\max}, \frac{1}{R} \right) \cdot (3.0 \cdot (1.0 - w_i)^{1.5} + 0.3), \quad (10)$$

where w_i is the normalized sensitivity weight and R is the target compression ratio. More sensitive tensors receive tighter budgets, clamped to $[10^{-4}, 0.05]$.

3.3 Dynamic Method Selection

The `DynamicIntelligenceSelector` scores all registered methods via:

$$S(m, t) = \underbrace{T_{\text{score}}(m)}_{\text{tier priority}} \times \underbrace{M(m, t)}_{\text{profile match}} \times \underbrace{P(m, t)}_{\text{predicted performance}} \times \underbrace{H(m, t)}_{\text{historical performance}} \quad (11)$$

where T_{score} assigns 10.0 for Tier 1 (decomposition, spectral, tensor network) down to 0.3 for Tier 5 (quantization). Methods are tried in score order; the first to meet the error budget while maximizing ratio is selected. Fallback chains ensure robustness: `block_int8` $\rightarrow$ `block_int4` $\rightarrow$ `passthrough`.

3.4 Method Taxonomy

The 80+ compression methods span 11 categories:

1. **Decomposition** (Tier 1): SVD, Tensor Train, CP, Tucker, Kronecker, CUR, Nyström, Butterfly, Monarch, H-matrix, Toeplitz, Hankel, block-diagonal.
2. **Spectral** (Tier 1): DCT 1D/2D, block DCT, FWHT, wavelets (Haar, Daubechies, Symlet, scattering), FFT, NTT, Givens, Chebyshev, Winograd, random Hadamard.
3. **Structural** (Tier 2): Low-rank, NMF, BlockSparse, Circulant, Vandermonde, Cauchy, HSS, BSS, structured 2:4 sparsity, SparseGPT, Wanda pruning.
4. **Entropy Coding** (Tier 3): Huffman, ANS, tANS, arithmetic, LZ77, DEFLATE, BWT+MTF, predictive coding.
5. **Functional** (Tier 1): MLP Mixer, FNet, Performer, LinearAttention, SIREN, neural ODE, symbolic regression, Fisher information.
6. **Physics-Inspired** (Tier 2): Vlasov mean-field, gyrokinetic, plasma oscillation, Debye shielding, density matrix, quantum state, quantum entanglement, topological data analysis.
7. **Tensor Network** (Tier 1): MERA, PEPS, QTT, DMRG, tensor ring, TT-cross, matrix product operator (MPO), quantum circuit simulation.
8. **Novel** (Tier 1): QuantumPlasmaFusion, HDC compression, HolographicReduceRank, Born machine, topological defects, time crystal, superposition decoding.
9. **Hybrid/Cascade** (Tier 4): 2/3/4-stage cascades combining transform $\rightarrow$ quantize $\rightarrow$ sparsify $\rightarrow$ entropy code in arbitrary order.
10. **Lossless** (Tier 3): Zstd, zlib, LZ4, RANS.
11. **Quantization** (Tier 5): Block INT8/4, Hadamard INT8/4, NF4, K-means, binary/ternary, GPTQ, AWQ, SqueezeLLM, lattice, product, mixed-precision.

For extreme compression ratios ($> 100 : 1$), the `MultiplicativeStackingEngine` chains methods (e.g., DCT + Tensor Train + rANS), and the ultra-compression path uses SVD with auto-tuned rank via the elbow method on the singular value spectrum.


3.5 Compression Format

Each `CompressedTensor` stores: raw compressed bytes, method identifier, parameters (block sizes, ranks, thresholds, quantization scales), quality metrics (relative error, SNR, PSNR, cosine similarity, compression ratio), and quality grade (S/A/B/C/D/F). Grades correspond to error thresholds: S $< 0.02\%$, A $< 0.1\%$, B $< 0.5\%$, C $< 1\%$, D $< 5\%$, F $\geq 5\%$.

4 SpectralStream Format (SSF)

The SSF binary format (v2/v3) is a page-aligned container with cryptographic integrity verification.

4.1 Binary Layout



4.2 Header Structure

The 256-byte header (packed as `<4sIIIIQQQQQQ184s`) contains: magic bytes (SSF02), format version (2 or 3), flags, tensor count, file offsets for index, metadata, and tensor data, redundant header offset (at page 1 for recovery), and footer offset. Space for 184 zero-filled spare bytes ensures future extensibility.

4.3 Tensor Index

Each variable-length `TensorIndexEntry` encodes: name (UTF-8), shape, dtype (F32/F16/BF16/INT8/INT4/U8), compression method ID (mapped via a 79-method dictionary), JSON-encoded parameters, data offset/size, SHA-256 checksum of the compressed block, and five quality doubles (relative error, SNR, PSNR, cosine similarity, ratio).

4.4 Reader and Writer

`SSFReader` supports memory-mapped (mmap) zero-copy reads for large files, LRU tensor decompression cache, row-chunk reads for 2D tensors (without full decompression), cryptographic

checksum verification on every read, and subset extraction via `SSFWriter`. `SSFWriter` writes page-aligned (4096-byte) sections in a single pass, computes file-level SHA-256 checksums in the footer, and supports streaming (one-tensor-at-a-time) mode.

4.5 Format Bridges

`SSFCConverter` converts from GGUF (with GGML dequantization), HuggingFace safetensors, and NumPy dictionaries. The `_compress_via_engine` bridge maps method IDs to engine compression methods for on-the-fly decompression during read.

5 Hierarchical KV Cache

The KV cache system (~3000 lines) is a production-grade subsystem with four components: core data structures, manager, compressor, and eviction policies.

5.1 Cache Architecture

The `KVCacheManager` orchestrates per-layer `OrderedDict` caches with 40+ configuration parameters controlling compression method, quantize bits, Hadamard rotation, spectral keep fraction, wavelet decomposition level, SVD/TT rank, codebook size, and cache size limit (default 4 GB).

5.2 Adaptive Compression Selection

Compression method is selected based on cache pressure (ratio of current memory to limit):

$$\text{method} = \begin{cases} \text{fwht_int4} + \text{residual VQ}, & p > 0.8 \text{ (critical)}, \\ \text{fwht_int8} + \text{DCT sparse}, & 0.6 < p \leq 0.8 \text{ (high)}, \\ \text{Hadamard} + \text{spectral} + \text{wavelet}, & 0.4 < p \leq 0.6 \text{ (moderate)}, \\ \text{configured method}, & p \leq 0.4 \text{ (low)}. \end{cases} \quad (13)$$

5.3 Eviction Policies

Thirteen eviction policies are implemented:

1. **SpectralEviction** (default): composite score of spectral entropy, Landau-Zener co-

herence decay (half-life 1000 steps), recency, and access frequency.

2. **H2OEviction**: bottom percentile by attention score [15].
3. **SlidingWindow**: oldest position evicted.
4. **StreamingLLM**: protects first 4 tokens [16].
5. **ResonanceEviction**: physics-inspired, tracks frequency-domain resonance via FFT of score history.
6. **EntropyEviction**: minimum spectral entropy.
7. **PredictiveEviction**: linear regression on access history.
8. **ClusteringEviction**: K-means clustering of key vectors.
9. **TopologicalEviction**: eigendecomposition of key similarity matrix.
10. **ReinforcementLearning**: 4-arm bandit with ϵ -greedy ($\epsilon = 0.1$).
11. **HybridEviction**: voting ensemble of 4 policies.
12. **StalenessAware**: exponential time decay.
13. **EntropyGradient**: minimum entropy rate of change.

5.4 Cache Compressor

Thirty compression method implementations operate on individual KV entries, organized into families: transform-domain quantization (FWHT-int8/4, Hadamard+Lloyd-Max, DCT-sparse, FFT, wavelets, Chebyshev), low-rank (SVD, tensor train), vector quantization (K-means, low)Q, product quantization, 3-stage residual VQ), temporal/predictive (delta encoding, AR prediction order 4, cross-layer delta), physics-inspired (holographic HRR binding, Vlasov distribution, quantum state angle encoding, time crystal periodic approximation), adaptive (context-aware method selection based on spectral entropy, heavy hitter protection), and lattice (E8 Gosset lattice quantization).

5.5 Additional Features

Auto-tuning every 50 steps adjusts quantization bits based on hit rate; progressive compression re-compresses with stronger methods at > 70% pressure; SSD tiering offloads cold entries beyond 50% memory usage; cache coherence monitoring computes key-vector similarity fragmentation metrics; and simulated annealing optimizes eviction temperature.

6 Inference Pipeline

The `InferencePipeline` (670 lines) provides production-grade CPU inference with a three-tier model loading fallback chain: SSF reader → `ModelLoader` (binary parser) → `SafeTensorsLoader` (HF adapter with bfloat16-to-float32 conversion).

6.1 Model Architecture Support

The pipeline supports Gemma 4 (E2B: 35 layers, 1536 hidden, 8/1 GQA; E4B: 42 layers, 2560 hidden, 8/2 GQA) with full architectural features: grouped query attention with RoPE ($\theta = 10^6$), sliding window (512 tokens, 5 of 6 layers), shared KV layers (last 20), embedding scaling by $\sqrt{d_{\text{model}}}$, weight tying, GELU-tanh activation, Gemma4RMSNorm, and attention/logit softcapping. Configuration parameters extend to 131,072 token context and 262,144 vocabulary.

6.2 Forward Pass

For each transformer layer: RMSNorm → QKV projections → RoPE rotation → multi-head GQA with sliding window and attention softcap → output projection with residual → FFN RMSNorm → gated SiLU/GELU FFN with residual. The Gemma4FFN uses the tanh approximation:

$$\text{GELU}_{\text{tanh}}(x) = 0.5x \left(1 + \tanh \left(\sqrt{2/\pi} (x + 0.044715x^3) \right) \right) \quad (14)$$

6.3 Speculative Decoding

The `SpeculativeDecoder` implements draft-verify decoding [17] with $\gamma = 4$ draft tokens per verification pass, adaptive acceptance rate

tracking, and best-of-N draft candidate selection from the HDC draft engine.

6.4 Generation

Sampling pipeline: temperature scaling → top-K filtering via `argpartition` → softmax → top-P (nucleus) filtering → categorical sampling. Streaming generation yields tokens as produced. Perplexity measurement on random 1024-token windows is available.

6.5 Benchmark Suite

The `BenchmarkRunner` provides six benchmark categories: throughput (tokens/second at batch sizes 1/2/4 and sequence lengths 128/512/2048), per-method compression evaluation, quality distribution analysis, memory profiling, multi-configuration comparison grid search, and Pareto frontier identification for optimal (ratio, error) tradeoffs. Output formats include JSON, LaTeX tables, and human-readable summaries.

7 Certificate and Validation System

Professional-grade certificates provide auditable quality assurance for every compression operation.

7.1 Certificate Hierarchy

Three certificate types are generated:

1. **TensorCertificate**: per-tensor metrics—name, shape, original/compressed bytes, ratio, method, category, relative error, SNR, PSNR, cosine similarity, MSE, timing, quality grade (S/A/B/C/D/F).
2. **CompressionCertificate**: full model—architecture, parameters, overall ratio, weighted/avg/max error, avg SNR, grade and method distributions, per-method-grade breakdown, and industry comparison against FP16 (2×), INT8 (4×), INT4 (8×), NF4 (4×), GPTQ (8×), AWQ (8×), GGML Q4_0 (4.5×), GGML Q8_0 (2.5×), and SqueezeLLM (8×).

3. **ValidationCertificate**: SSF structural integrity—header, checksum, index verification, per-tensor decompression and comparison against originals.

7.2 Output Formats

All three certificate types generate four output formats: JSON (machine-readable), HTML (dark-themed with badges, progress bars, and gradient headers), Markdown (README-ready tables), and plain text (terminal display with box-drawing characters). The `CertificateBuilder` provides factory methods `from_compression_report()` and `from_compressed_tensors()`.

7.3 End-to-End Validation Pipeline

The `scripts/e2e_validation.py` script orchestrates the full pipeline: synthetic model generation (low-rank structure with controlled noise to simulate real compressible weights), engine compression round-trip, SSF format write/read round-trip, multi-format certificate generation, and threshold enforcement (default: avg error $\leq 1\%$, overall ratio $\geq 50 : 1$). Exit code 0 indicates all thresholds met; exit code 1 indicates breach with detailed diagnostics.

8 Performance and Characteristics

8.1 Compression Performance

The intelligence engine achieves extreme compression ratios through the synergistic combination of methods. Spectral methods (Tier 1) typically achieve 20–200 $\times$ on single tensors; tensor network methods (MERA, TT) reach 200–77,000 $\times$ on structured matrices; cascade methods compound these multiplicatively. The adaptive selector ensures that each tensor receives the method best suited to its specific structure.

8.2 Computational Efficiency

All operations are vectorized via NumPy. Key operations achieve: FWHT $O(N \log N)$ via butterfly; 2D DCT $O(N^2)$ via separable matrix multiply; truncated SVD $O(\min(mn^2, m^2n))$ via full SVD or $O(mnk)$ via randomized SVD [18]. The

pure-Python implementation with no C++ extensions dramatically simplifies deployment, auditing, and cross-platform compatibility, at the cost of being bound by NumPy’s BLAS backend for large tensor operations.

8.3 Memory Architecture

The system manages memory at multiple levels: tensor-level streaming (one tensor at a time), SSF-level memory-mapped access with LRU decompressed tensor cache (configurable GB), KV cache size limits with automatic SSD tiering, process memory tracking via psutil, and automatic garbage collection with GC-triggered streaming fallback.

8.4 Scalability

The framework handles models with up to hundreds of billions of parameters via streaming mode, 131,072-token contexts via the KV cache system, 262,144 vocabulary via HD-based speculative decoding that avoids full softmax, and up to 42-layer architectures with shared KV and sliding window attention.

9 Conclusion

SpectralStream provides a unified framework for neural network compression and efficient inference that spans six mathematical traditions, 80+ compression methods, a production-grade inference pipeline, a hierarchical KV cache with 30+ compression methods and 13 eviction policies, and a professional certificate system. By reversing the conventional quantization-first approach and instead prioritizing spectral, structural, and tensor-network decompositions, SpectralStream achieves compression ratios that are orders of magnitude beyond conventional quantization.

The framework’s pure-Python implementation with NumPy vectorization demonstrates that high-performance neural network compression does not require C++ extensions or CUDA kernels, making it uniquely accessible for research, auditing, and customization. All components—from mathematical primitives to the inference pipeline and certificate system—are designed as modular, composable building blocks that can be extended, replaced, or studied independently.

Future Directions

Planned developments include GPU-accelerated tensor operations via JAX/CuPy backends, on-device inference for mobile and edge deployment, adaptive quantization that dynamically adjusts precision during generation, and integration with additional model architectures beyond the current Gemma 4 support.

References

- [1] N. Ahmed, T. Natarajan, and K. R. Rao, “Discrete cosine transform,” *IEEE Trans. Computers*, vol. C-23, no. 1, pp. 90–93, 1974.
- [2] P. Kanerva, “Hyperdimensional computing: An introduction to computing in distributed representation with high-dimensional random vectors,” *Cognitive Computation*, vol. 1, no. 2, pp. 139–159, 2009.
- [3] T. A. Plate, *Holographic Reduced Representation: Distributed Representation for Cognitive Structures*. CSLI Publications, 2003.
- [4] A. A. Vlasov, “The vibrational properties of an electron gas,” *Soviet Physics Uspekhi*, vol. 10, no. 6, pp. 721–733, 1968.
- [5] R. Orús, “A practical introduction to tensor networks: Matrix product states and projected entangled pair states,” *Annals of Physics*, vol. 349, pp. 117–158, 2014.
- [6] G. Vidal, “Entanglement renormalization,” *Phys. Rev. Lett.*, vol. 99, no. 22, p. 220405, 2007.
- [7] A. Beck and M. Teboulle, “A fast iterative shrinkage-thresholding algorithm for linear inverse problems,” *SIAM J. Imaging Sciences*, vol. 2, no. 1, pp. 183–202, 2009.
- [8] J. Duda, K. Tahboub, N. J. Gadgil, and E. J. Delp, “The use of asymmetric numeral systems as an accurate replacement for Huffman coding,” in *Proc. Picture Coding Symposium*, 2015, pp. 65–69.
- [9] E. Frantar, S. Ashkboos, T. Hoefer, and D. Alistarh, “GPTQ: Accurate post-training quantization for generative pre-trained transformers,” in *Proc. ICLR*, 2023.
- [10] J. Lin *et al.*, “AWQ: Activation-aware weight quantization for LLM compression and acceleration,” in *Proc. MLSys*, 2024.
- [11] E. Frantar and D. Alistarh, “SparseGPT: Massive language models can be accurately pruned in one-shot,” in *Proc. ICML*, 2023.
- [12] M. Sun *et al.*, “Wanda: A simple and effective pruning approach for large language models,” in *Proc. NeurIPS*, 2023.
- [13] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” arXiv:1503.02531, 2015.
- [14] G. Gerganov, “ggml: Tensor library for machine learning,” <https://github.com/ggerganov/ggml>, 2023.
- [15] Z. Zhang *et al.*, “H2O: Heavy-hitter oracle for efficient generative inference of large language models,” in *Proc. NeurIPS*, 2023.
- [16] G. Xiao *et al.*, “StreamingLLM: Efficient streaming language models with attention sinks,” arXiv:2309.17453, 2023.
- [17] Y. Leviathan, M. Kalman, and Y. Matias, “Fast inference from transformers via speculative decoding,” in *Proc. ICML*, 2023.
- [18] N. Halko, P. G. Martinsson, and J. A. Tropp, “Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions,” *SIAM Review*, vol. 53, no. 2, pp. 217–288, 2011.